




Neural Network Models


for natural language processing

Neural Network Model

Neural Network Model




Neural Network Model



Distributed architecture
Connectionist basis

Neural Network Model



What is a model?

Neural Network Model



Useful representation

Often but not always
an abstraction

~~Neural Network~~ Model

Models

**For any word from vocabulary,
with a single guess**

- 1. predict its frequency**
- 2. predict the next word**

Loss/cost function

$$\lambda(y, \hat{y}) = (y - \hat{y})^2$$

quadratic loss

$$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y})$$

hinge loss

$$\log p(y \mid \theta)$$

log predictive density

Loss/cost function

$\lambda(y, \hat{y}) = (y - \hat{y})^2$ quadratic loss

$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y})$ hinge loss

$\log p(y \mid \theta)$ log predictive density

To be a good model, you need to know what you will be evaluated on!

Loss/cost function

$\lambda(y, \hat{y}) = (y - \hat{y})^2$ quadratic loss


$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y})$ hinge loss

$\log p(y \mid \theta)$ log predictive density

There is no free lunch

For any word v from vocabulary of size $|V|$,
with a single guess

1. predict its frequency
2. predict the next word




For any word v from vocabulary of size $|V|$,
with a single **predictor**

1. predict its frequency
2. predict the next word

For any word v from vocabulary of size $|V|$,
with a single **predictor**

1. predict its frequency
2. predict the next word




For any word v from vocabulary of size $|V|$,
with a **multiple predictors**

1. predict its frequency
2. predict the next word

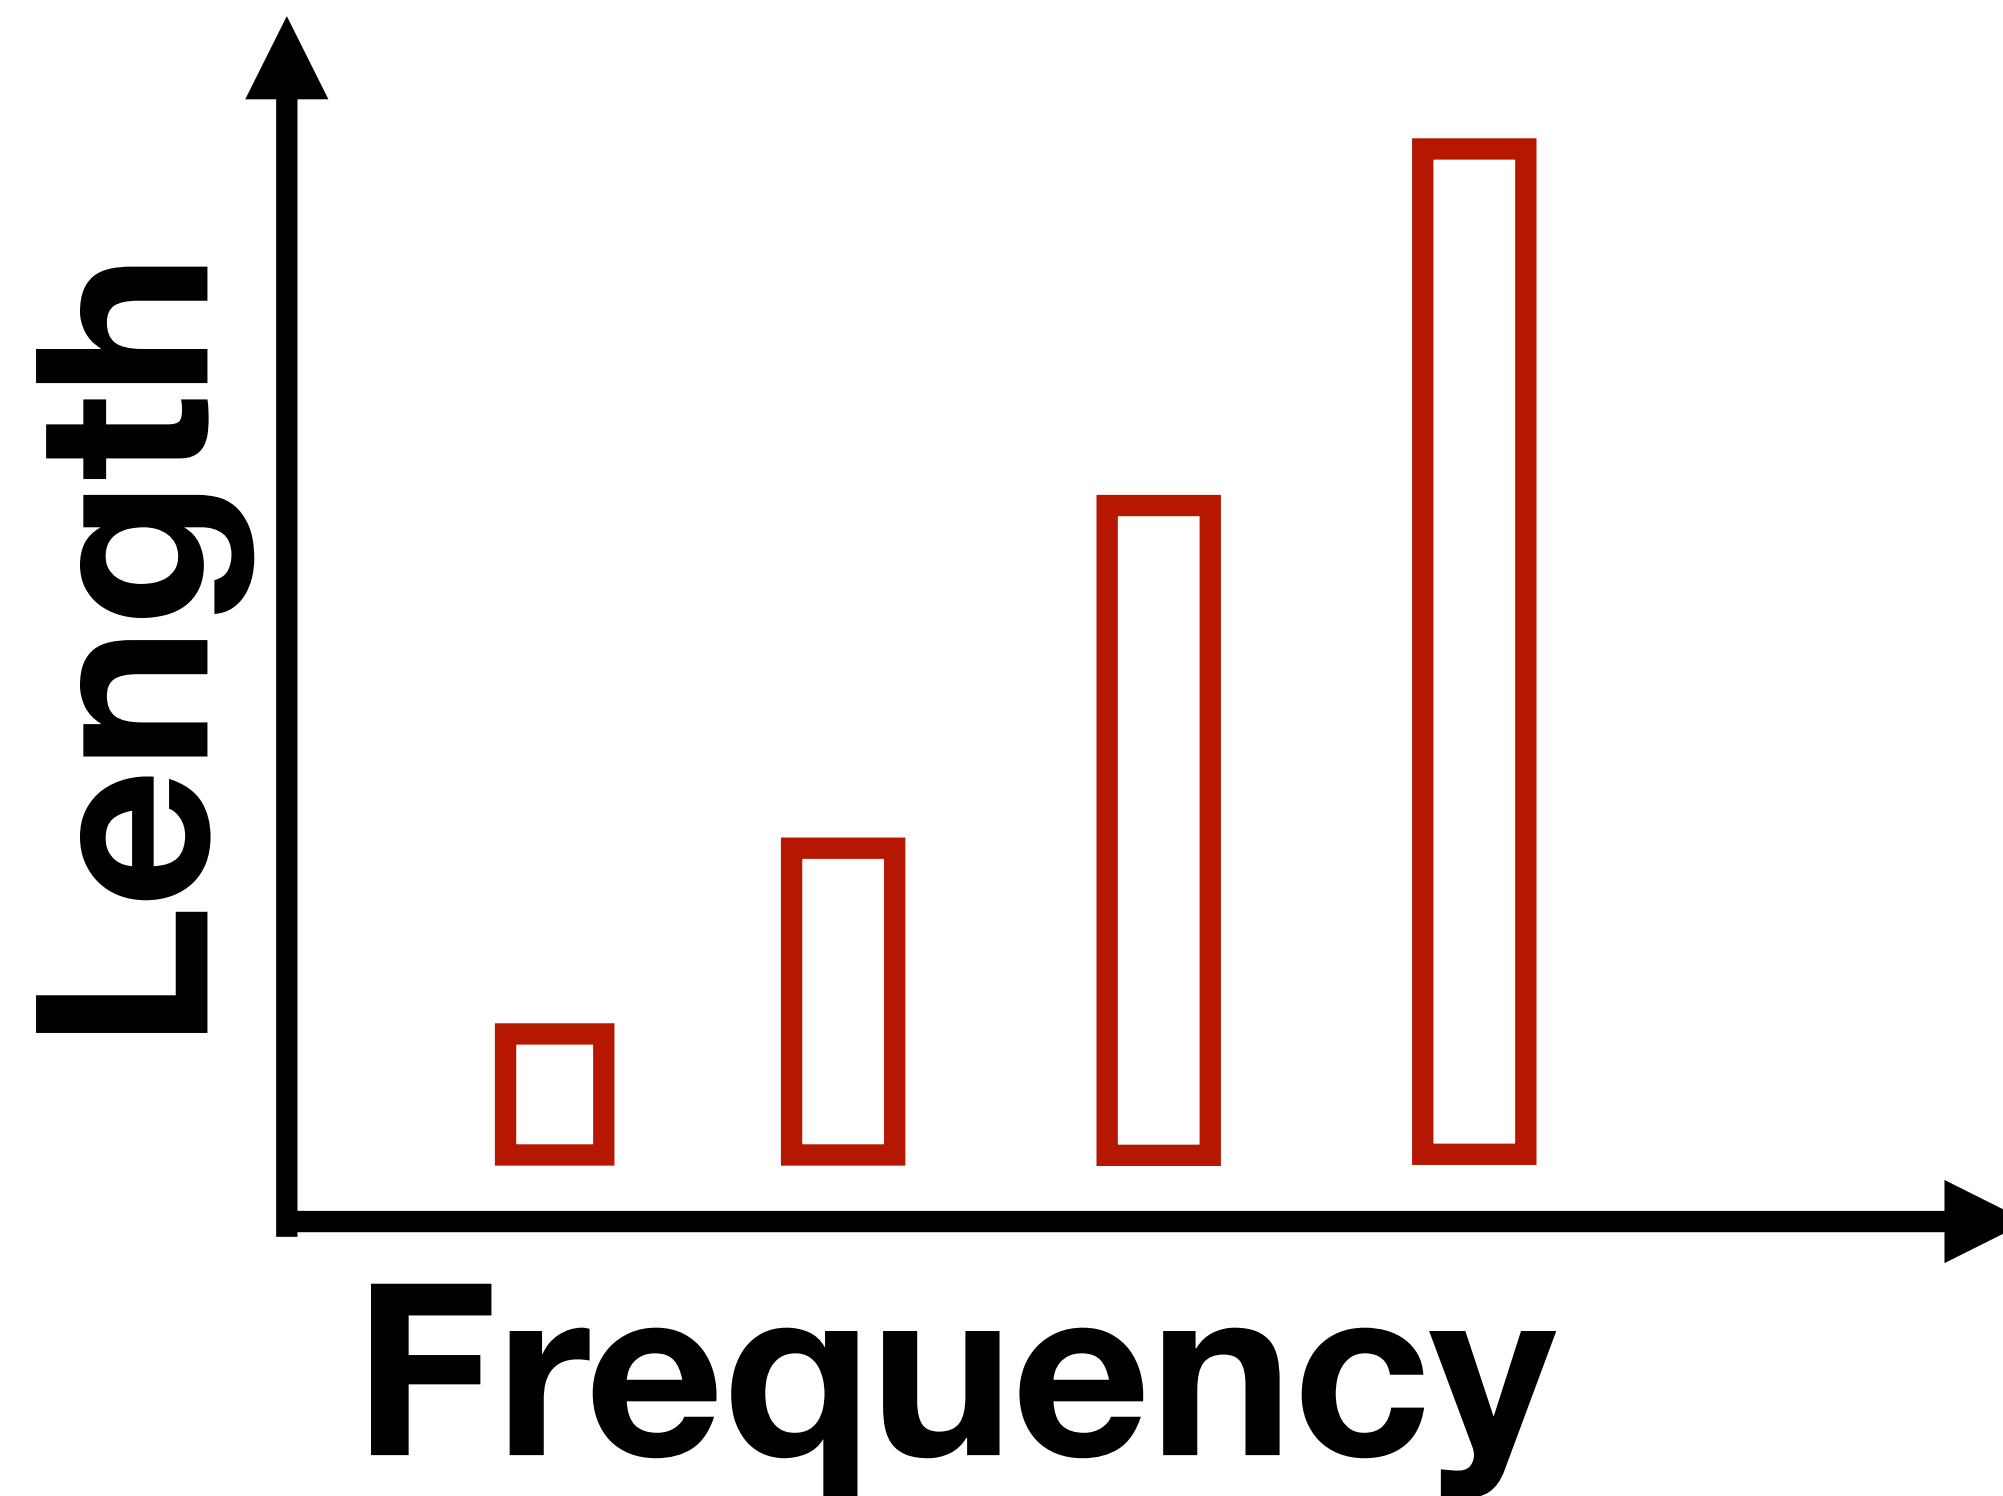
How can the multiple predictors be combined?

Linear models and non-linearity

$$\text{frequency}_i = \beta_0$$



$$\text{frequency}_i = \beta_0 + \beta_1 \text{length}_i$$




$$\text{frequency}_i = \beta_0 + \beta_1 \text{length}_i + \beta_2 \text{pos}_i$$

$$\text{frequency}_i = \beta_0 + \beta_1 \text{length}_i \times \beta_2 \text{pos}_i$$


...

You can keep making this more complex. The point is that this the outcome of a **linear combination of parameters**


next word_{*i*} = β_0




$$\text{pr}(\text{next word}_i) = \beta_0$$



$$\text{pr}(\text{next word}_i) = f(\beta_0)$$





Non-linearity, at last!

**Generalized linear models are still linear models
even though they use a non-linear transformation**

Generalized linear models are still linear models even though they use a non-linear transformation

They explicitly estimate the effect of one or more predictors on an outcome

Generalized linear models are still linear models even though they use a non-linear transformation

They explicitly estimate the effect of one or more predictors on an outcome

NNs scale up these ideas but are **non-linear and have **many parameters with no clear semantics** behind them**

Neural Network models

Non-linearity: Many phenomena are non-linear, so linearity is a potentially unnecessary constraint in the relationship between input and output

Parameters with no clear semantics: Automatically induced from data with no need to match architecture to phenomenon*

Many parameters: Can be an issue but doesn't need to be

Dense representations

Feed-forward NN

$$[x_1, \dots, x_{d_{in}}] \Rightarrow \text{NN}(\cdot) \Rightarrow [y_1, \dots, y_{d_{out}}]$$


Feed-forward NN

Output
layer


Hidden
layer

Hidden
layer

Input layer



Feed-forward NN



$$\sigma(\sum_1^3 w_i k_i + b) = g(x)$$


Feed-forward NN

Output
layer

Hidden
layer

Hidden
layer

Input layer



$$g^2(g^1(xW^1 + b^1)W^2 + b^2)$$

$$g^1(xW^1 + b^1)$$

$$x \in \mathbb{R}^{d_{in}}$$

Decisions for the modeller/engineer

1. Number of dimensions

Decisions for the modeller/engineer

1. Number of dimensions

2. Number of layers

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity
6. Loss function

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity
6. Loss function
7. Training regime
(e.g., stochastic gradient descent + flavor;
batching; drop-out)