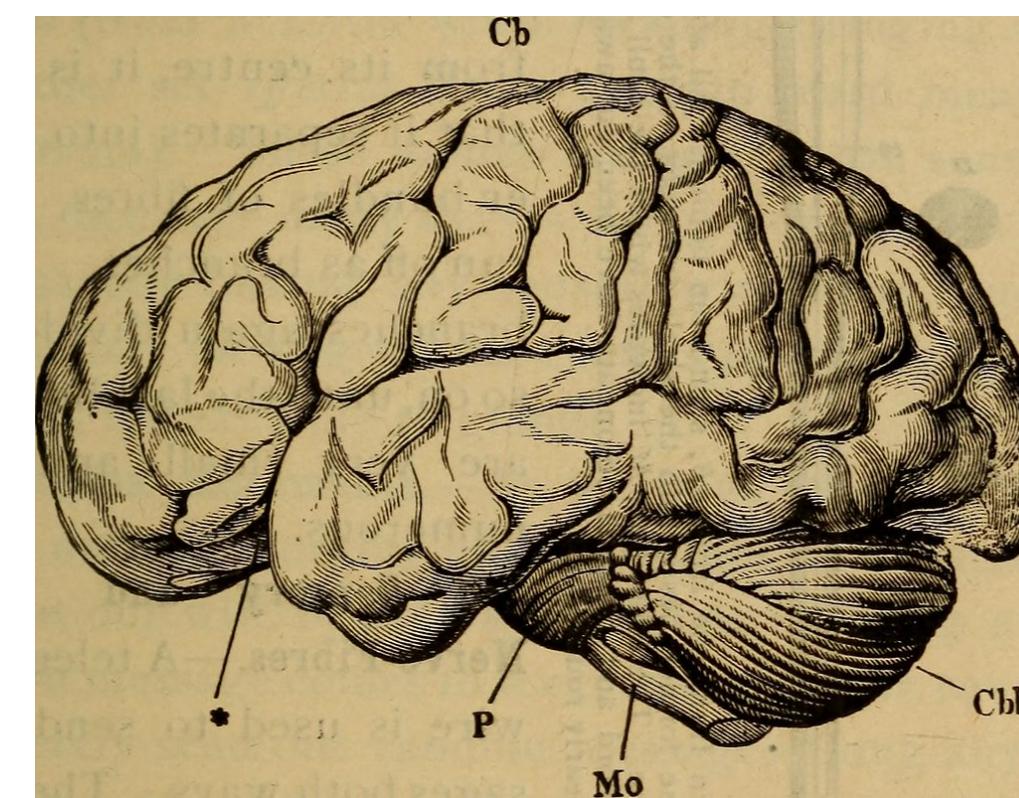
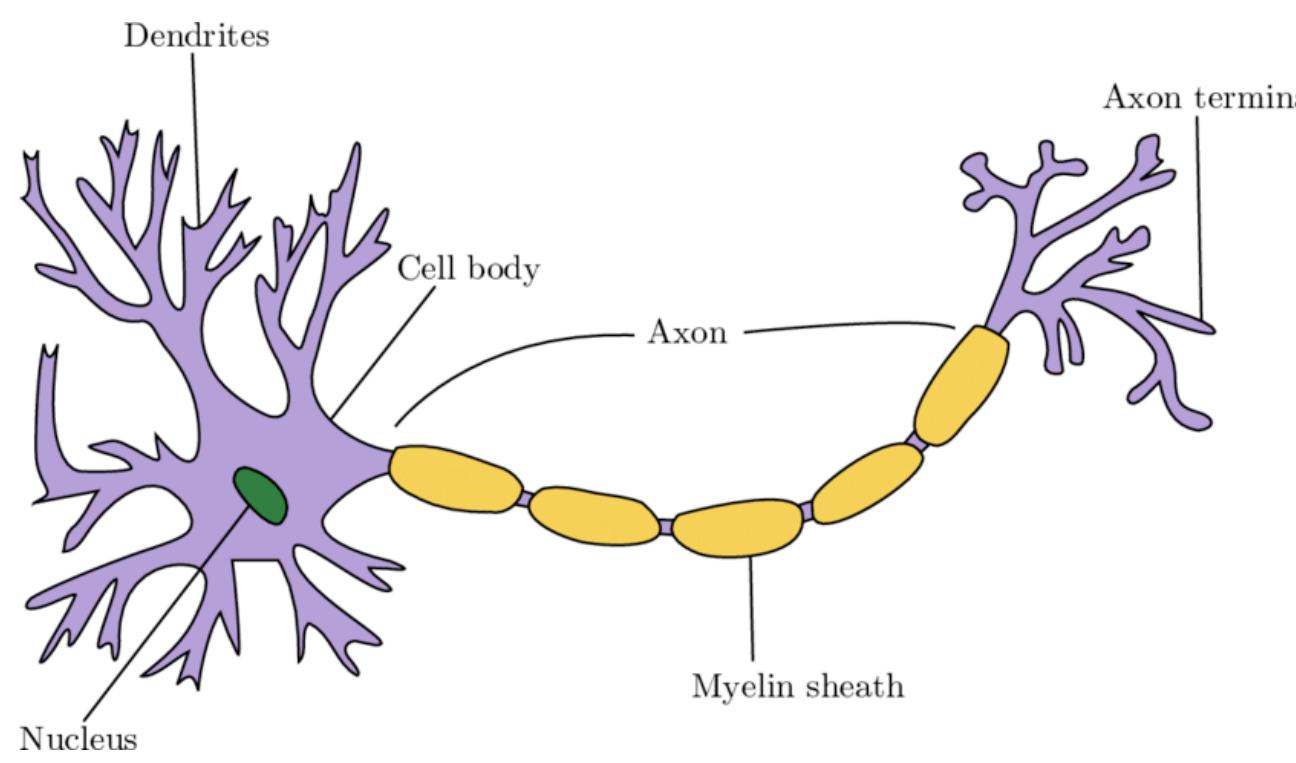
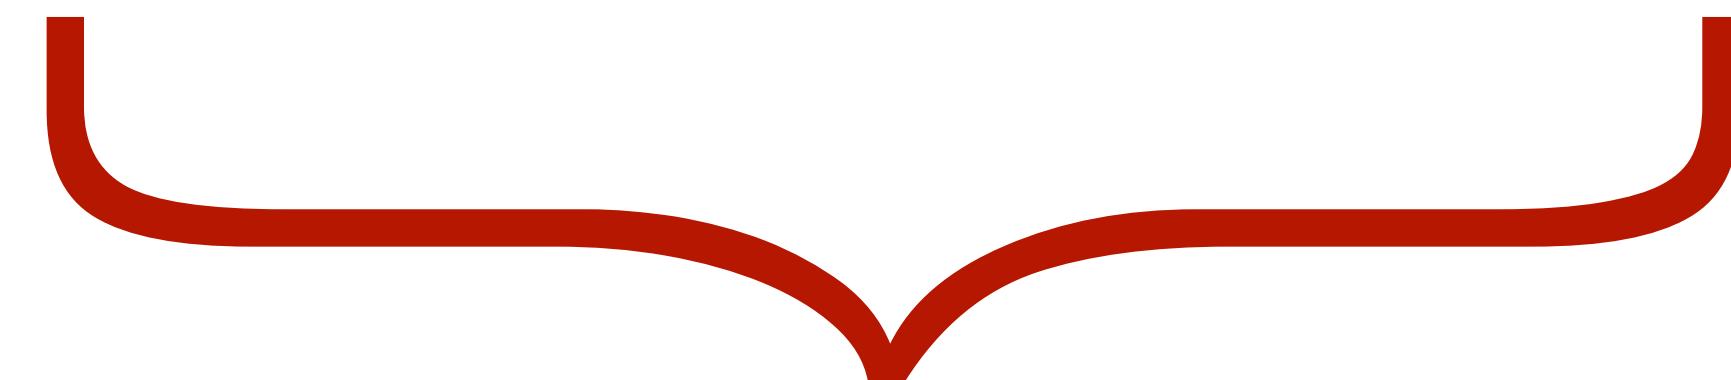
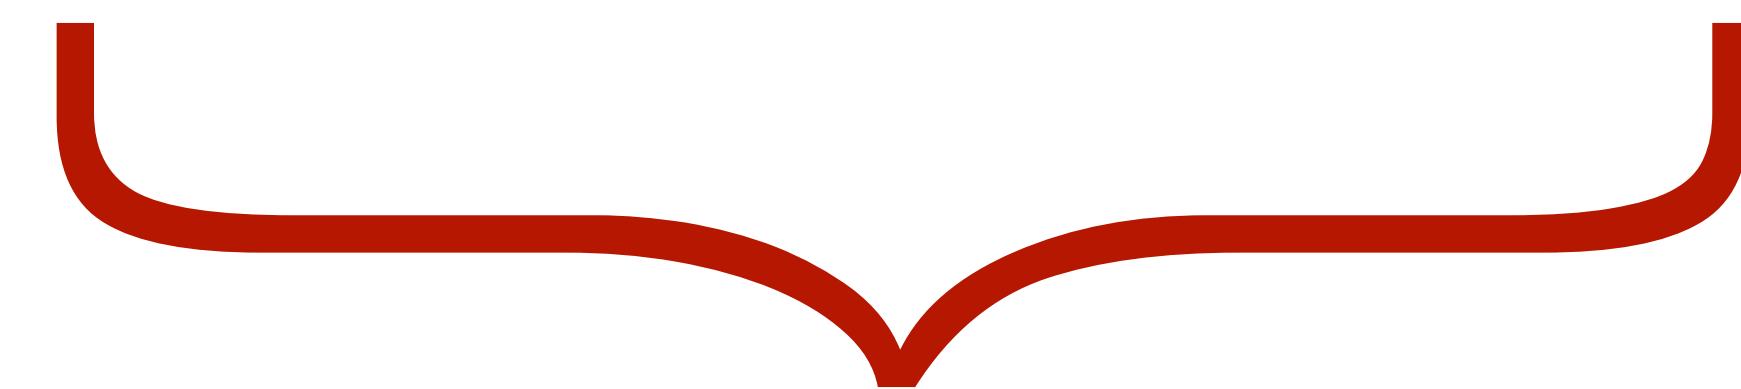


Neural Network Model

Neural Network Model

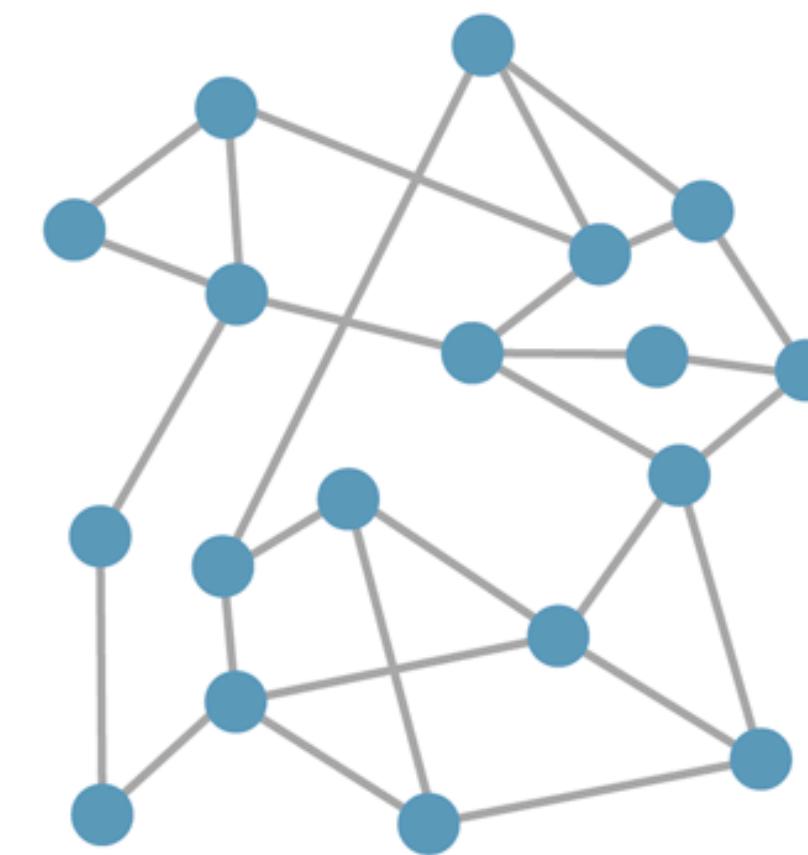


Neural Network Model



Distributed architecture
Connectionist basis

Neural Network Model



What is a model?

Neural Network Model



Useful representation

Often but not always
an abstraction



Ne**X**ral Network Model

Models

**For any word from vocabulary,
with a single guess**

- 1. predict its frequency**

- 2. predict the next word**

Loss/cost function

$$\lambda(y, \hat{y}) = (y - \hat{y})^2 \quad \text{quadratic loss}$$

$$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y}) \quad \text{hinge loss}$$

$$\log p(y \mid \theta) \quad \text{log predictive density}$$

Loss/cost function

$$\lambda(y, \hat{y}) = (y - \hat{y})^2 \quad \text{quadratic loss}$$

$$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y}) \quad \text{hinge loss}$$

$$\log p(y | \theta) \quad \text{log predictive density}$$

To be a good model, you need to know
what you will be evaluated on!

Loss/cost function

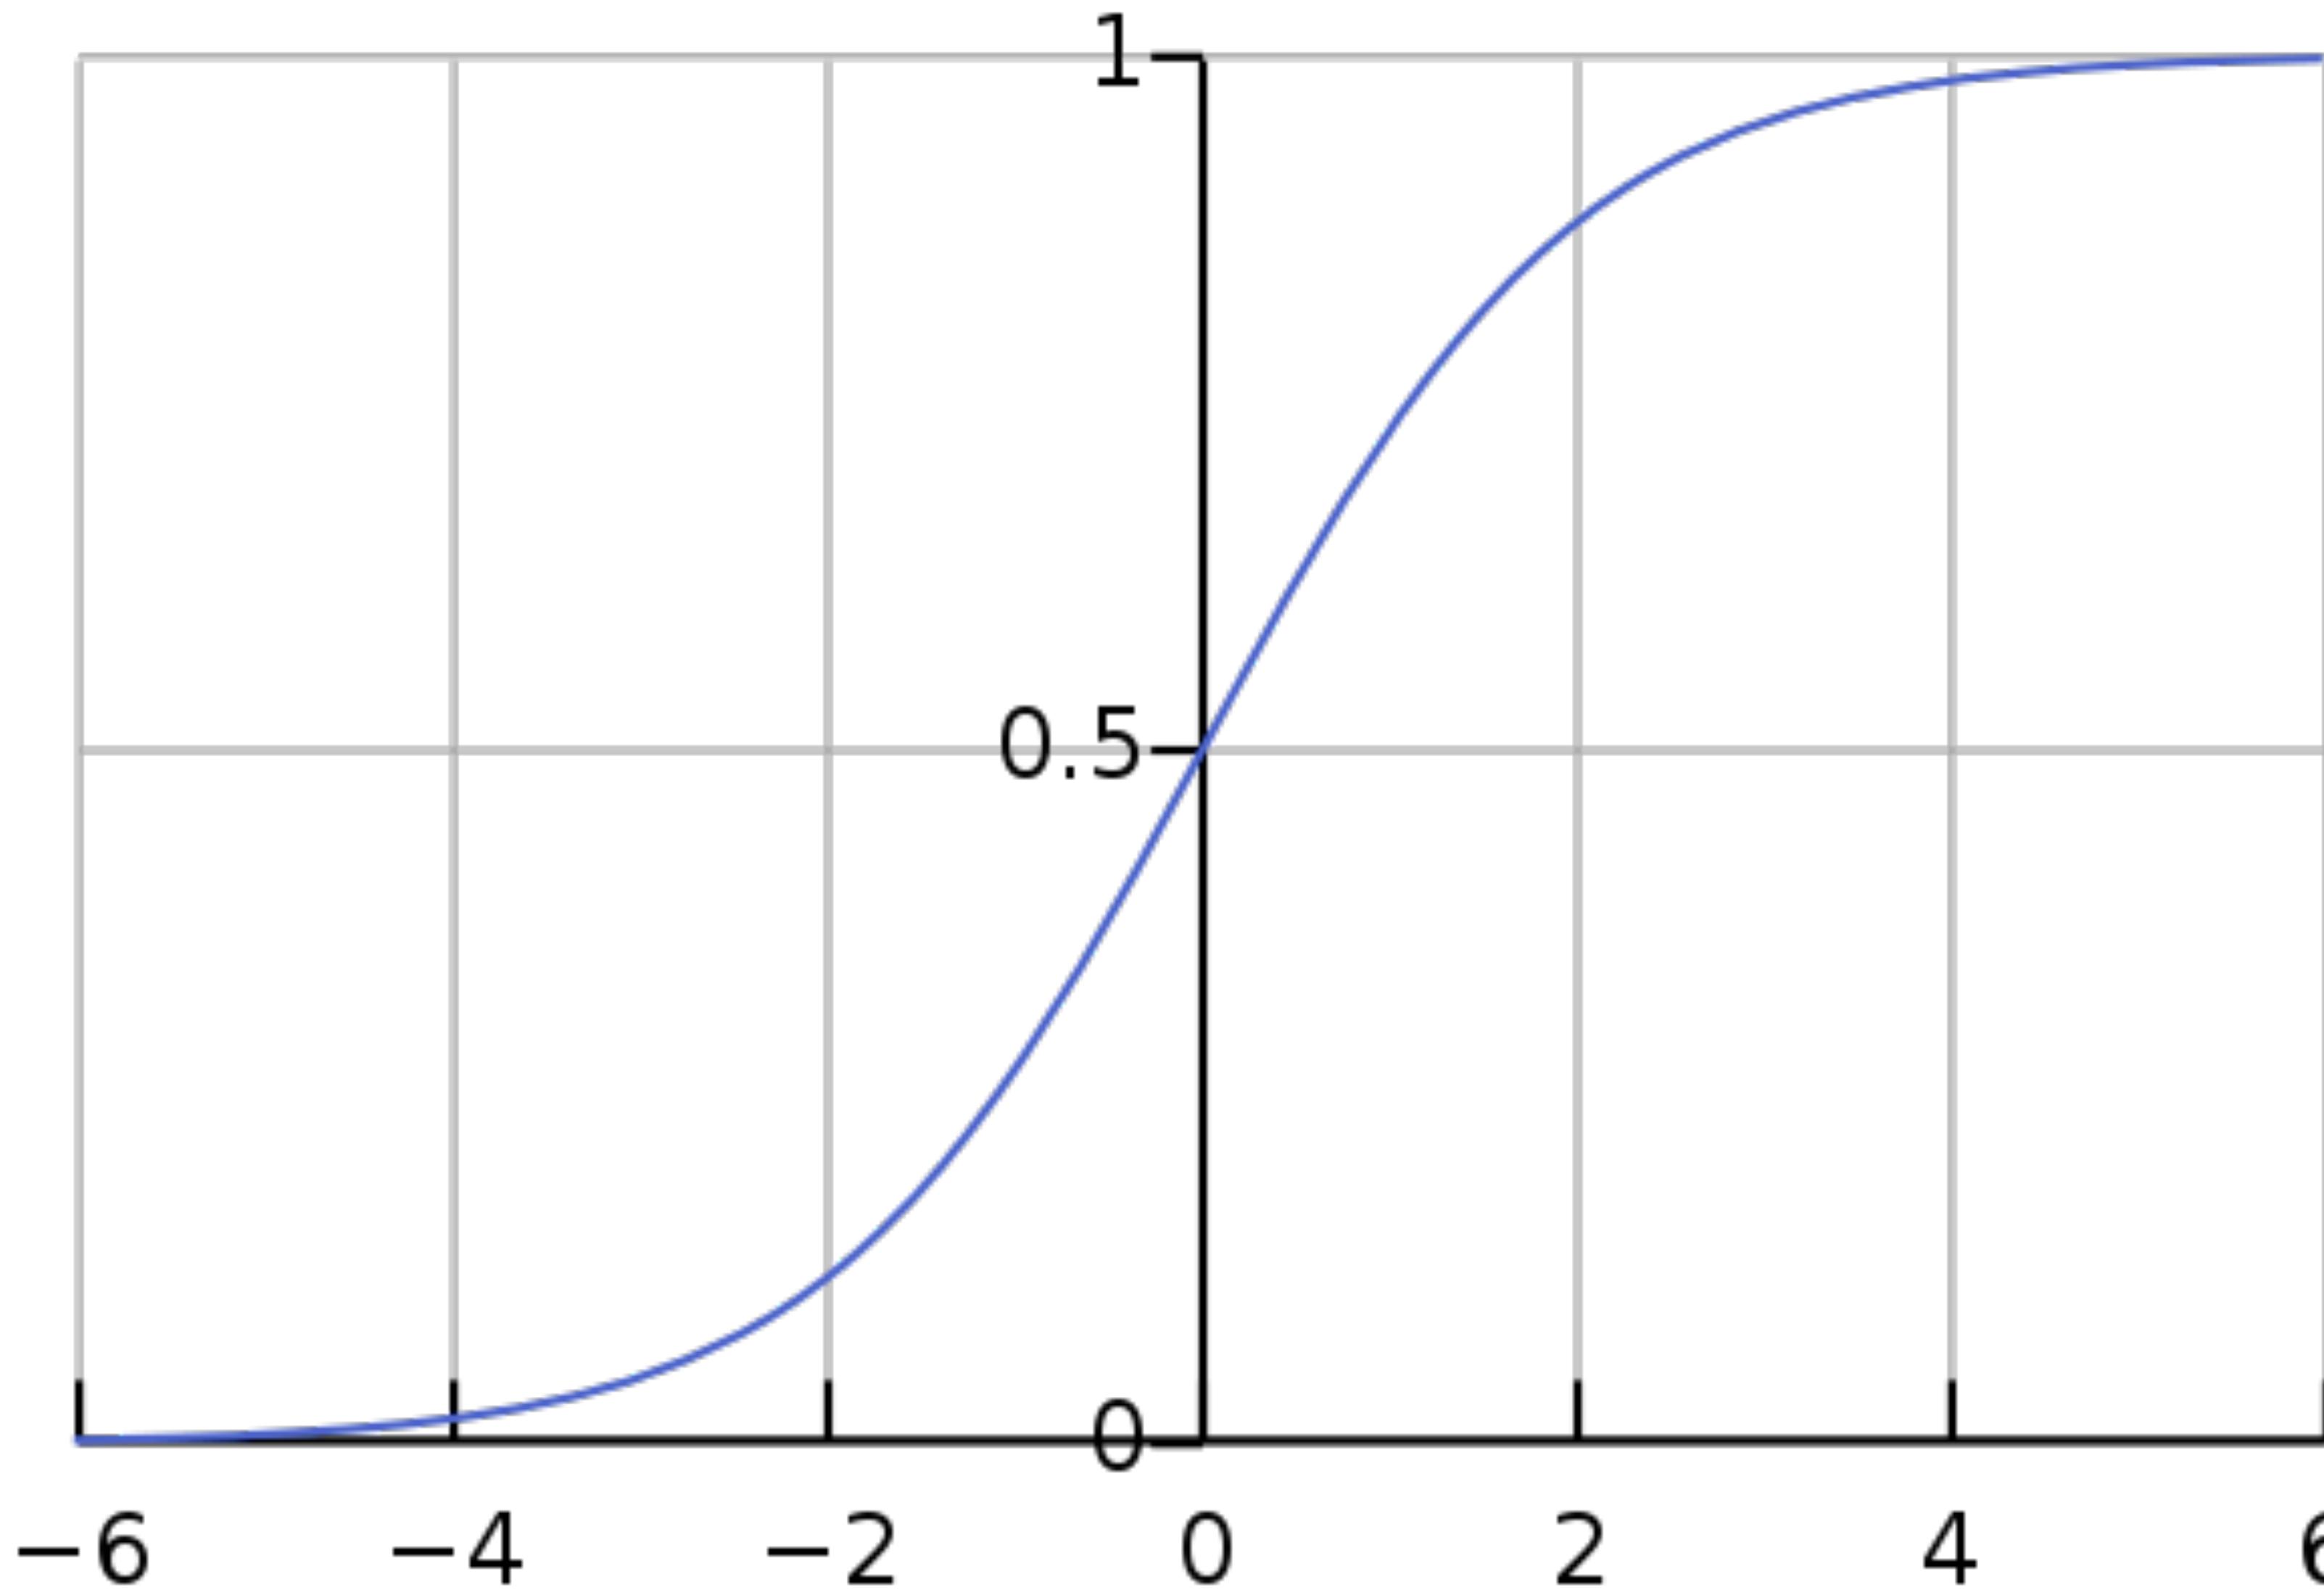
$$\lambda(y, \hat{y}) = (y - \hat{y})^2 \quad \text{quadratic loss}$$

$$\lambda(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y}) \quad \text{hinge loss}$$

$$\log p(y | \theta) \quad \text{log predictive density}$$

There is no free lunch

Linear models and non-linearity



Non-linearity, at last!

**Generalized linear models are still linear models
even though they use a non-linear transformation**

Generalized linear models are still linear models even though they use a non-linear transformation

They explicitly estimate the effect of one or more predictors on an outcome

Generalized linear models are still linear models even though they use a non-linear transformation

They explicitly estimate the effect of one or more predictors on an outcome

NNs scale up these ideas but are non-linear and have many parameters with no clear semantics behind them

Neural Network models

Non-linearity: Many phenomena are non-linear, so linearity is a potentially unnecessary constraint in the relationship between input and output

Parameters with no clear semantics: Automatically induced from data with no need to match architecture to phenomenon*

Many parameters: Can be an issue but doesn't need to be

Dense representations

Feed-forward NN

$$[x_1, \dots, x_{d_{in}}] \Rightarrow \text{NN}(\cdot) \Rightarrow [y_1, \dots, y_{d_{out}}]$$

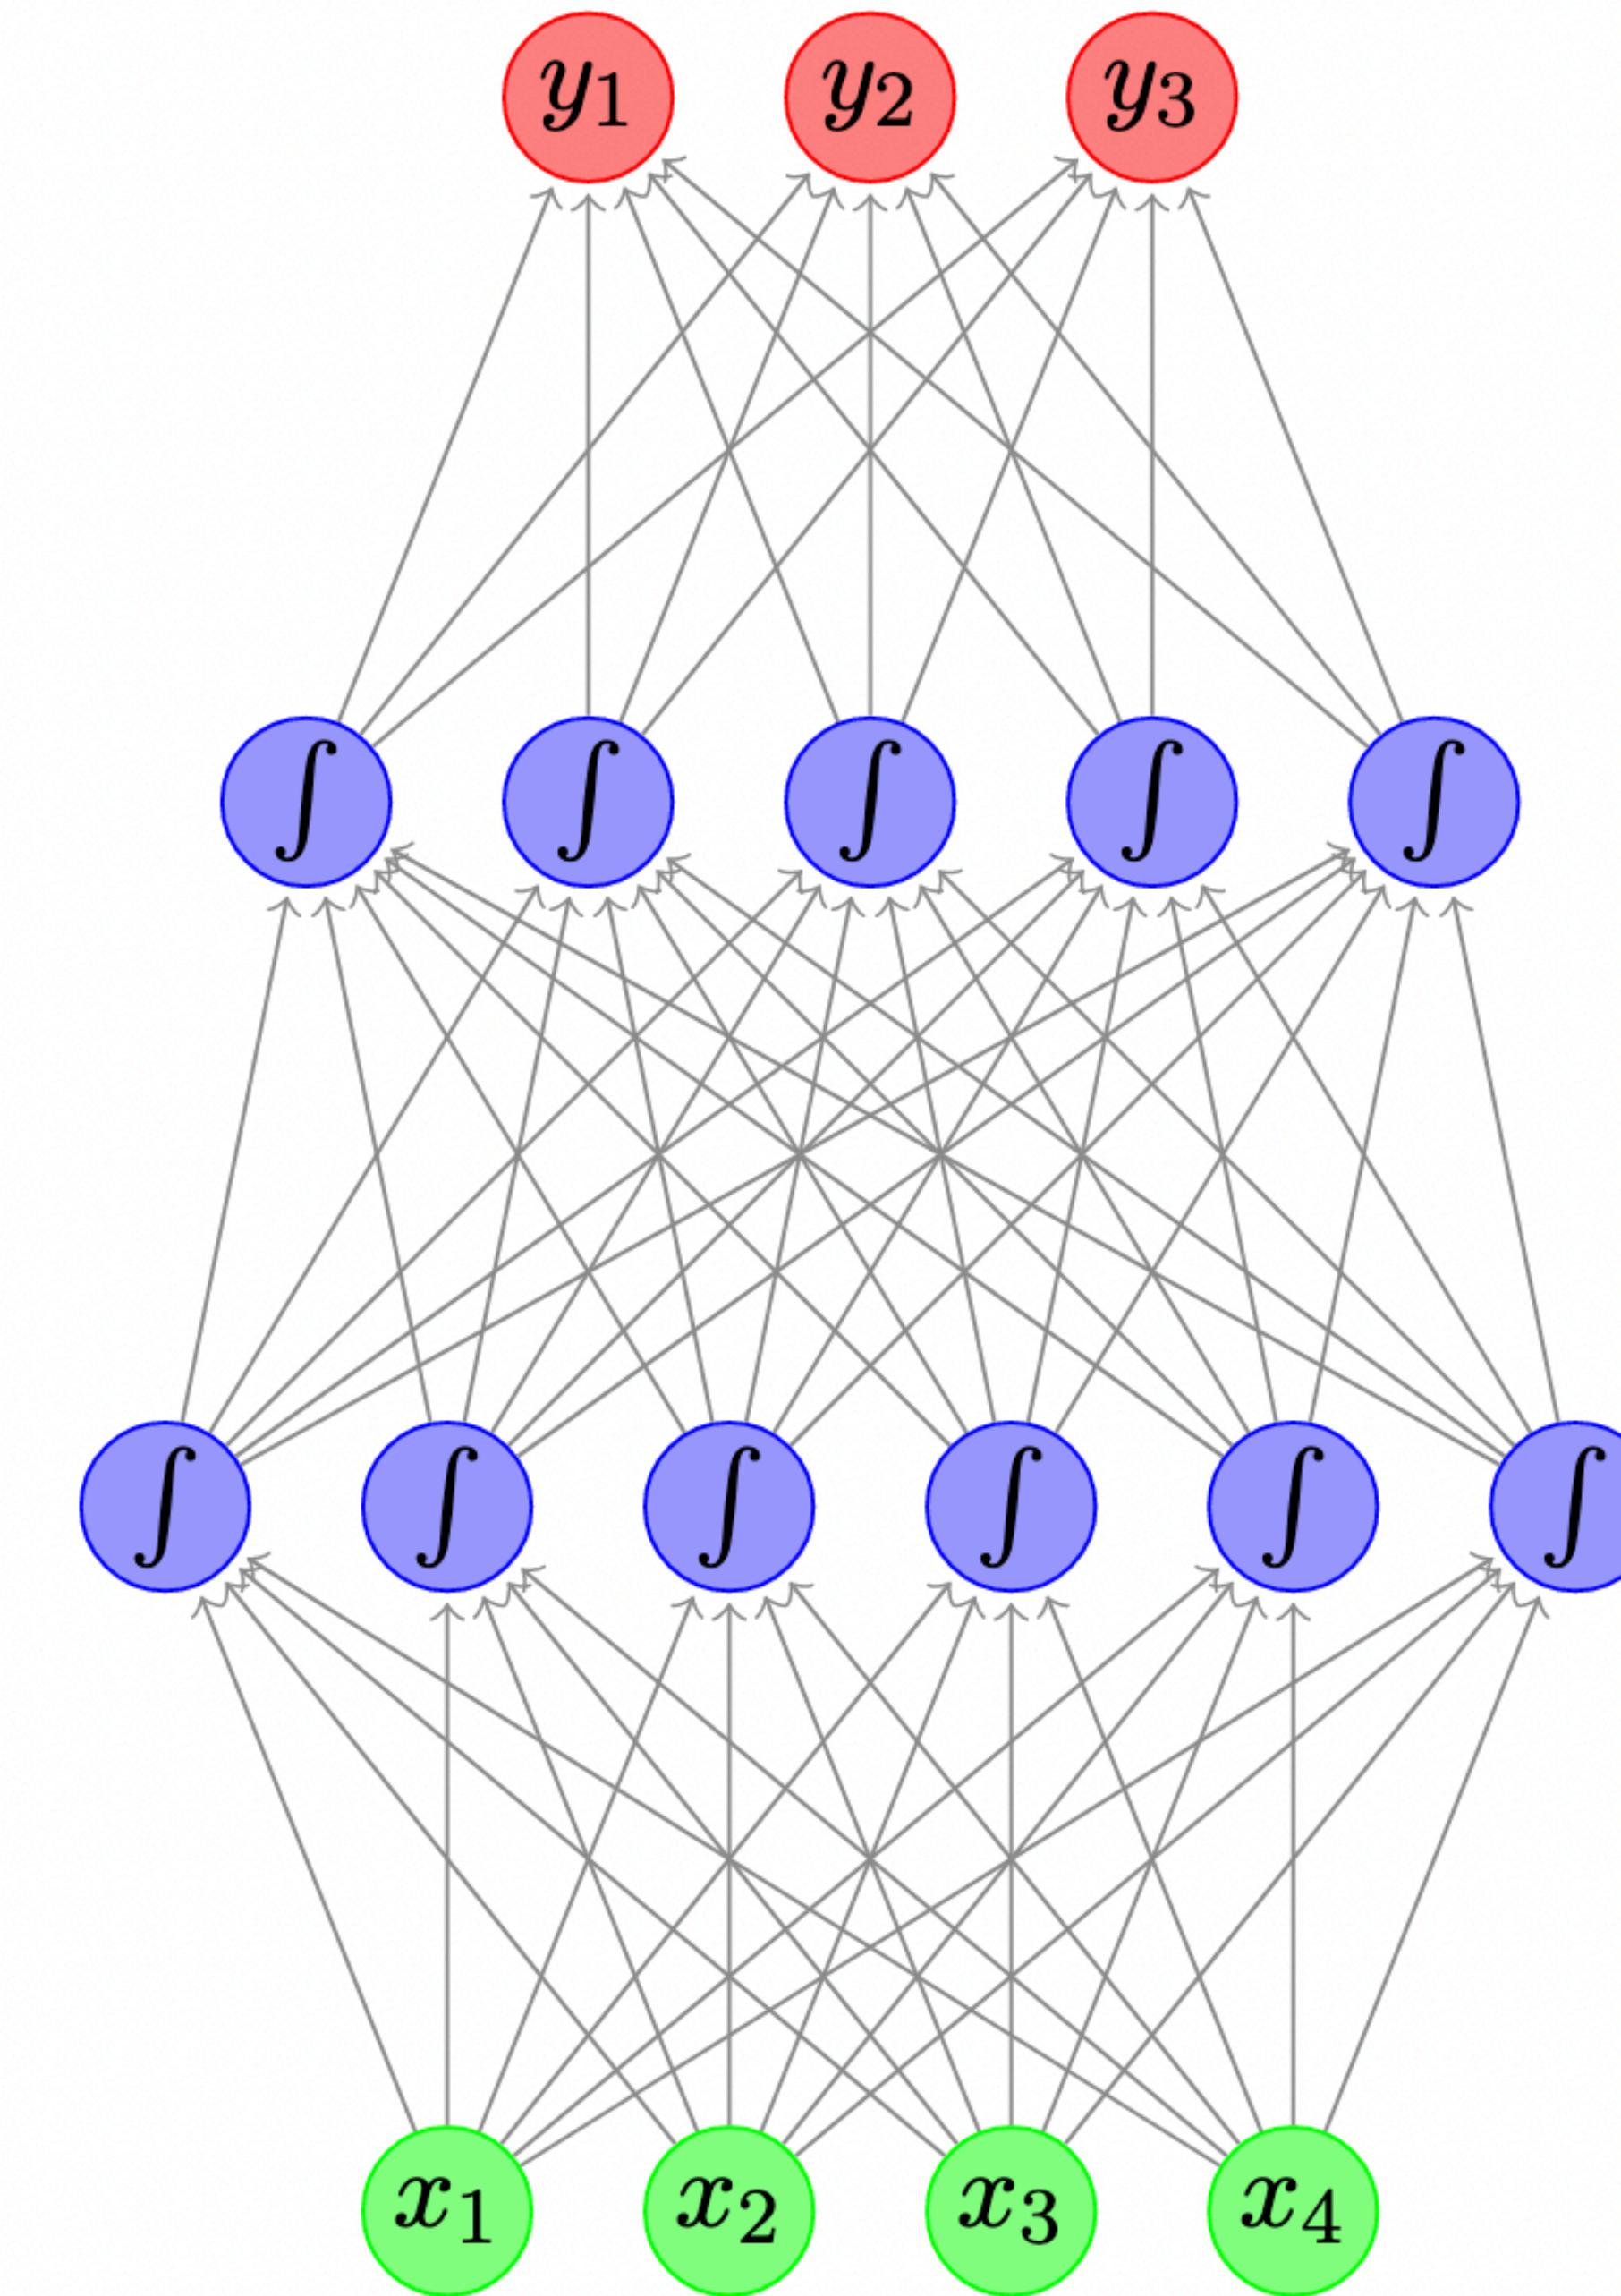
Feed-forward NN

Output
layer

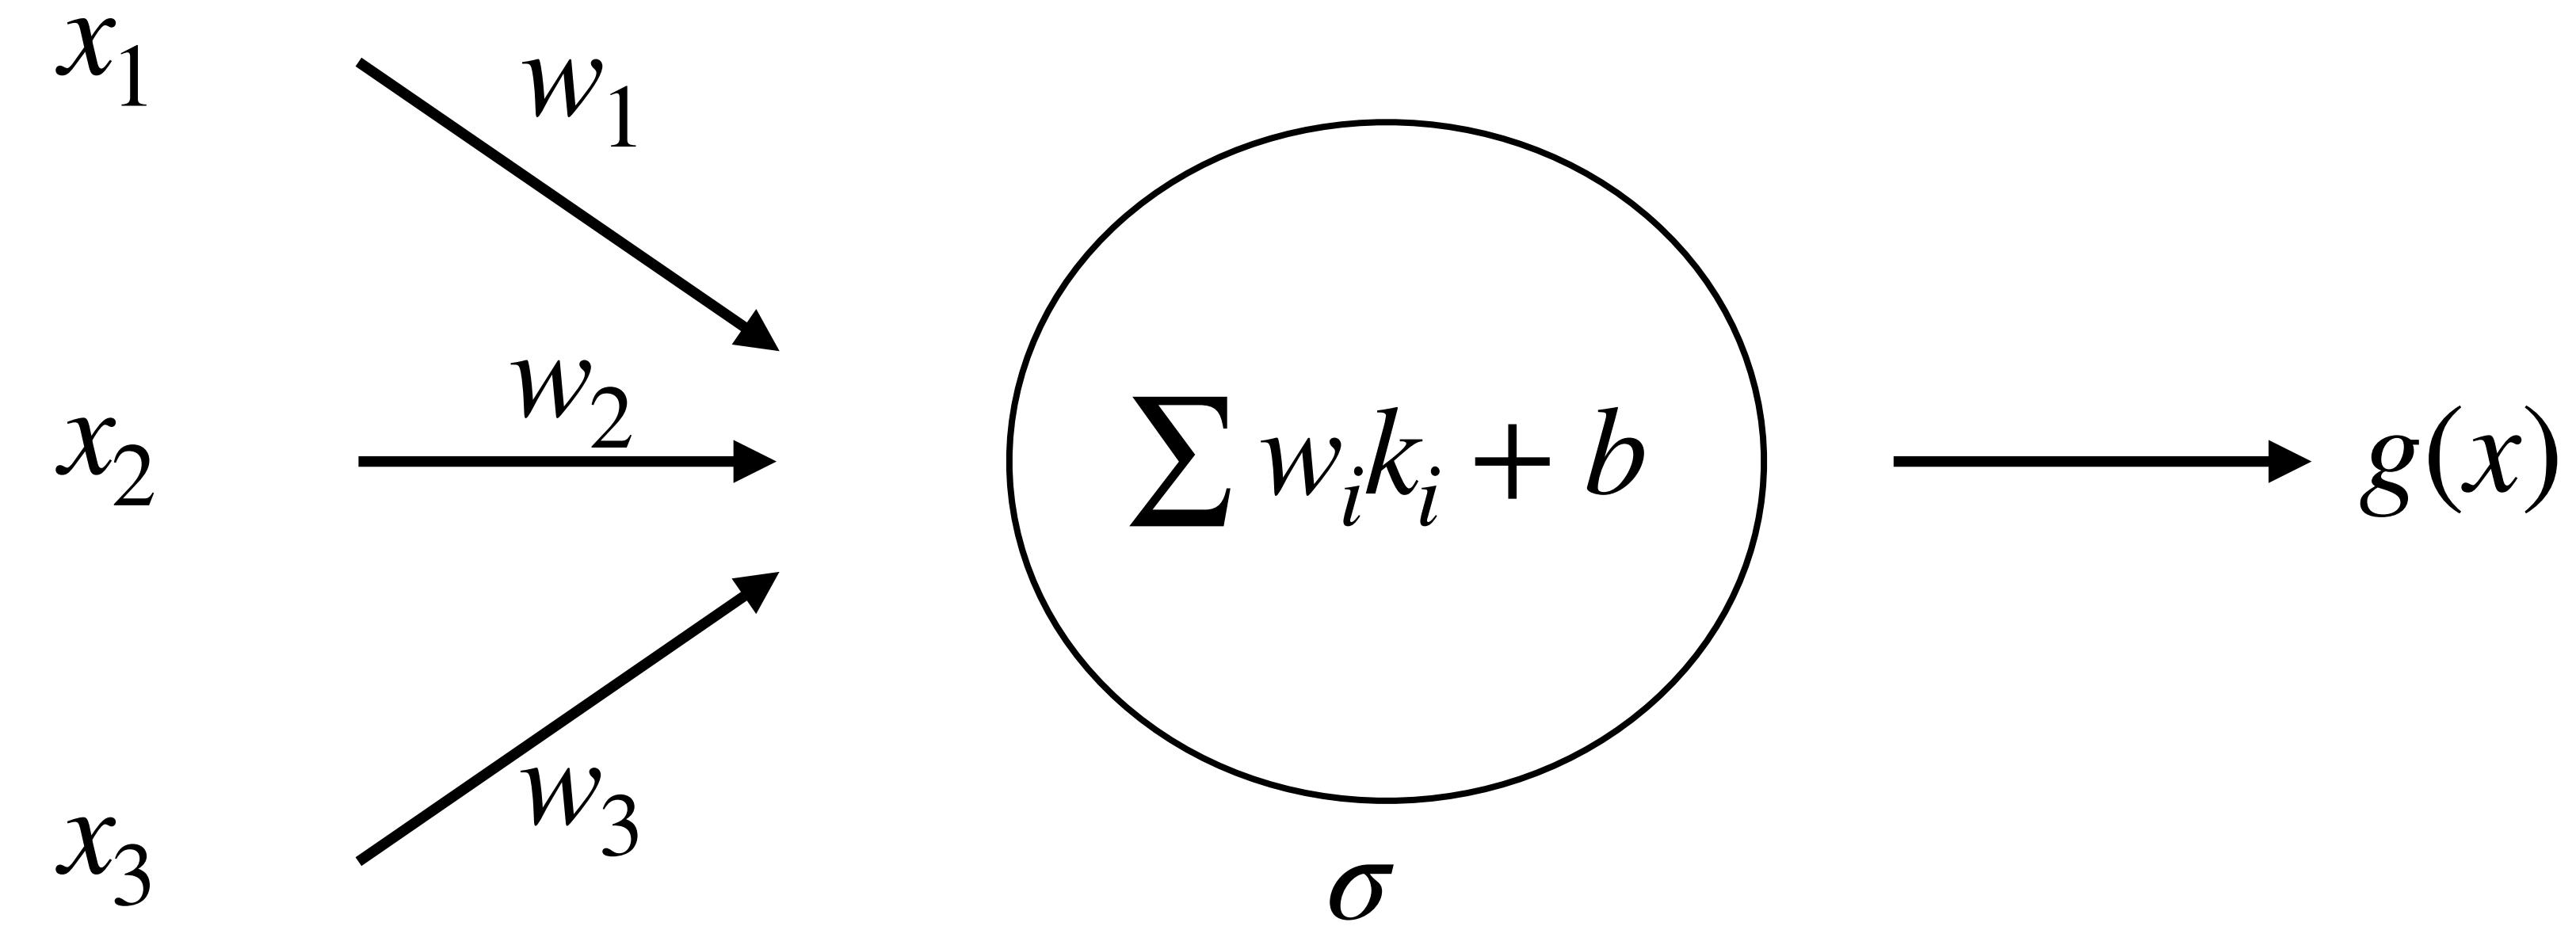
Hidden
layer

Hidden
layer

Input layer



Feed-forward NN



$$\sigma(\sum_1^3 w_i k_i + b) = g(x)$$

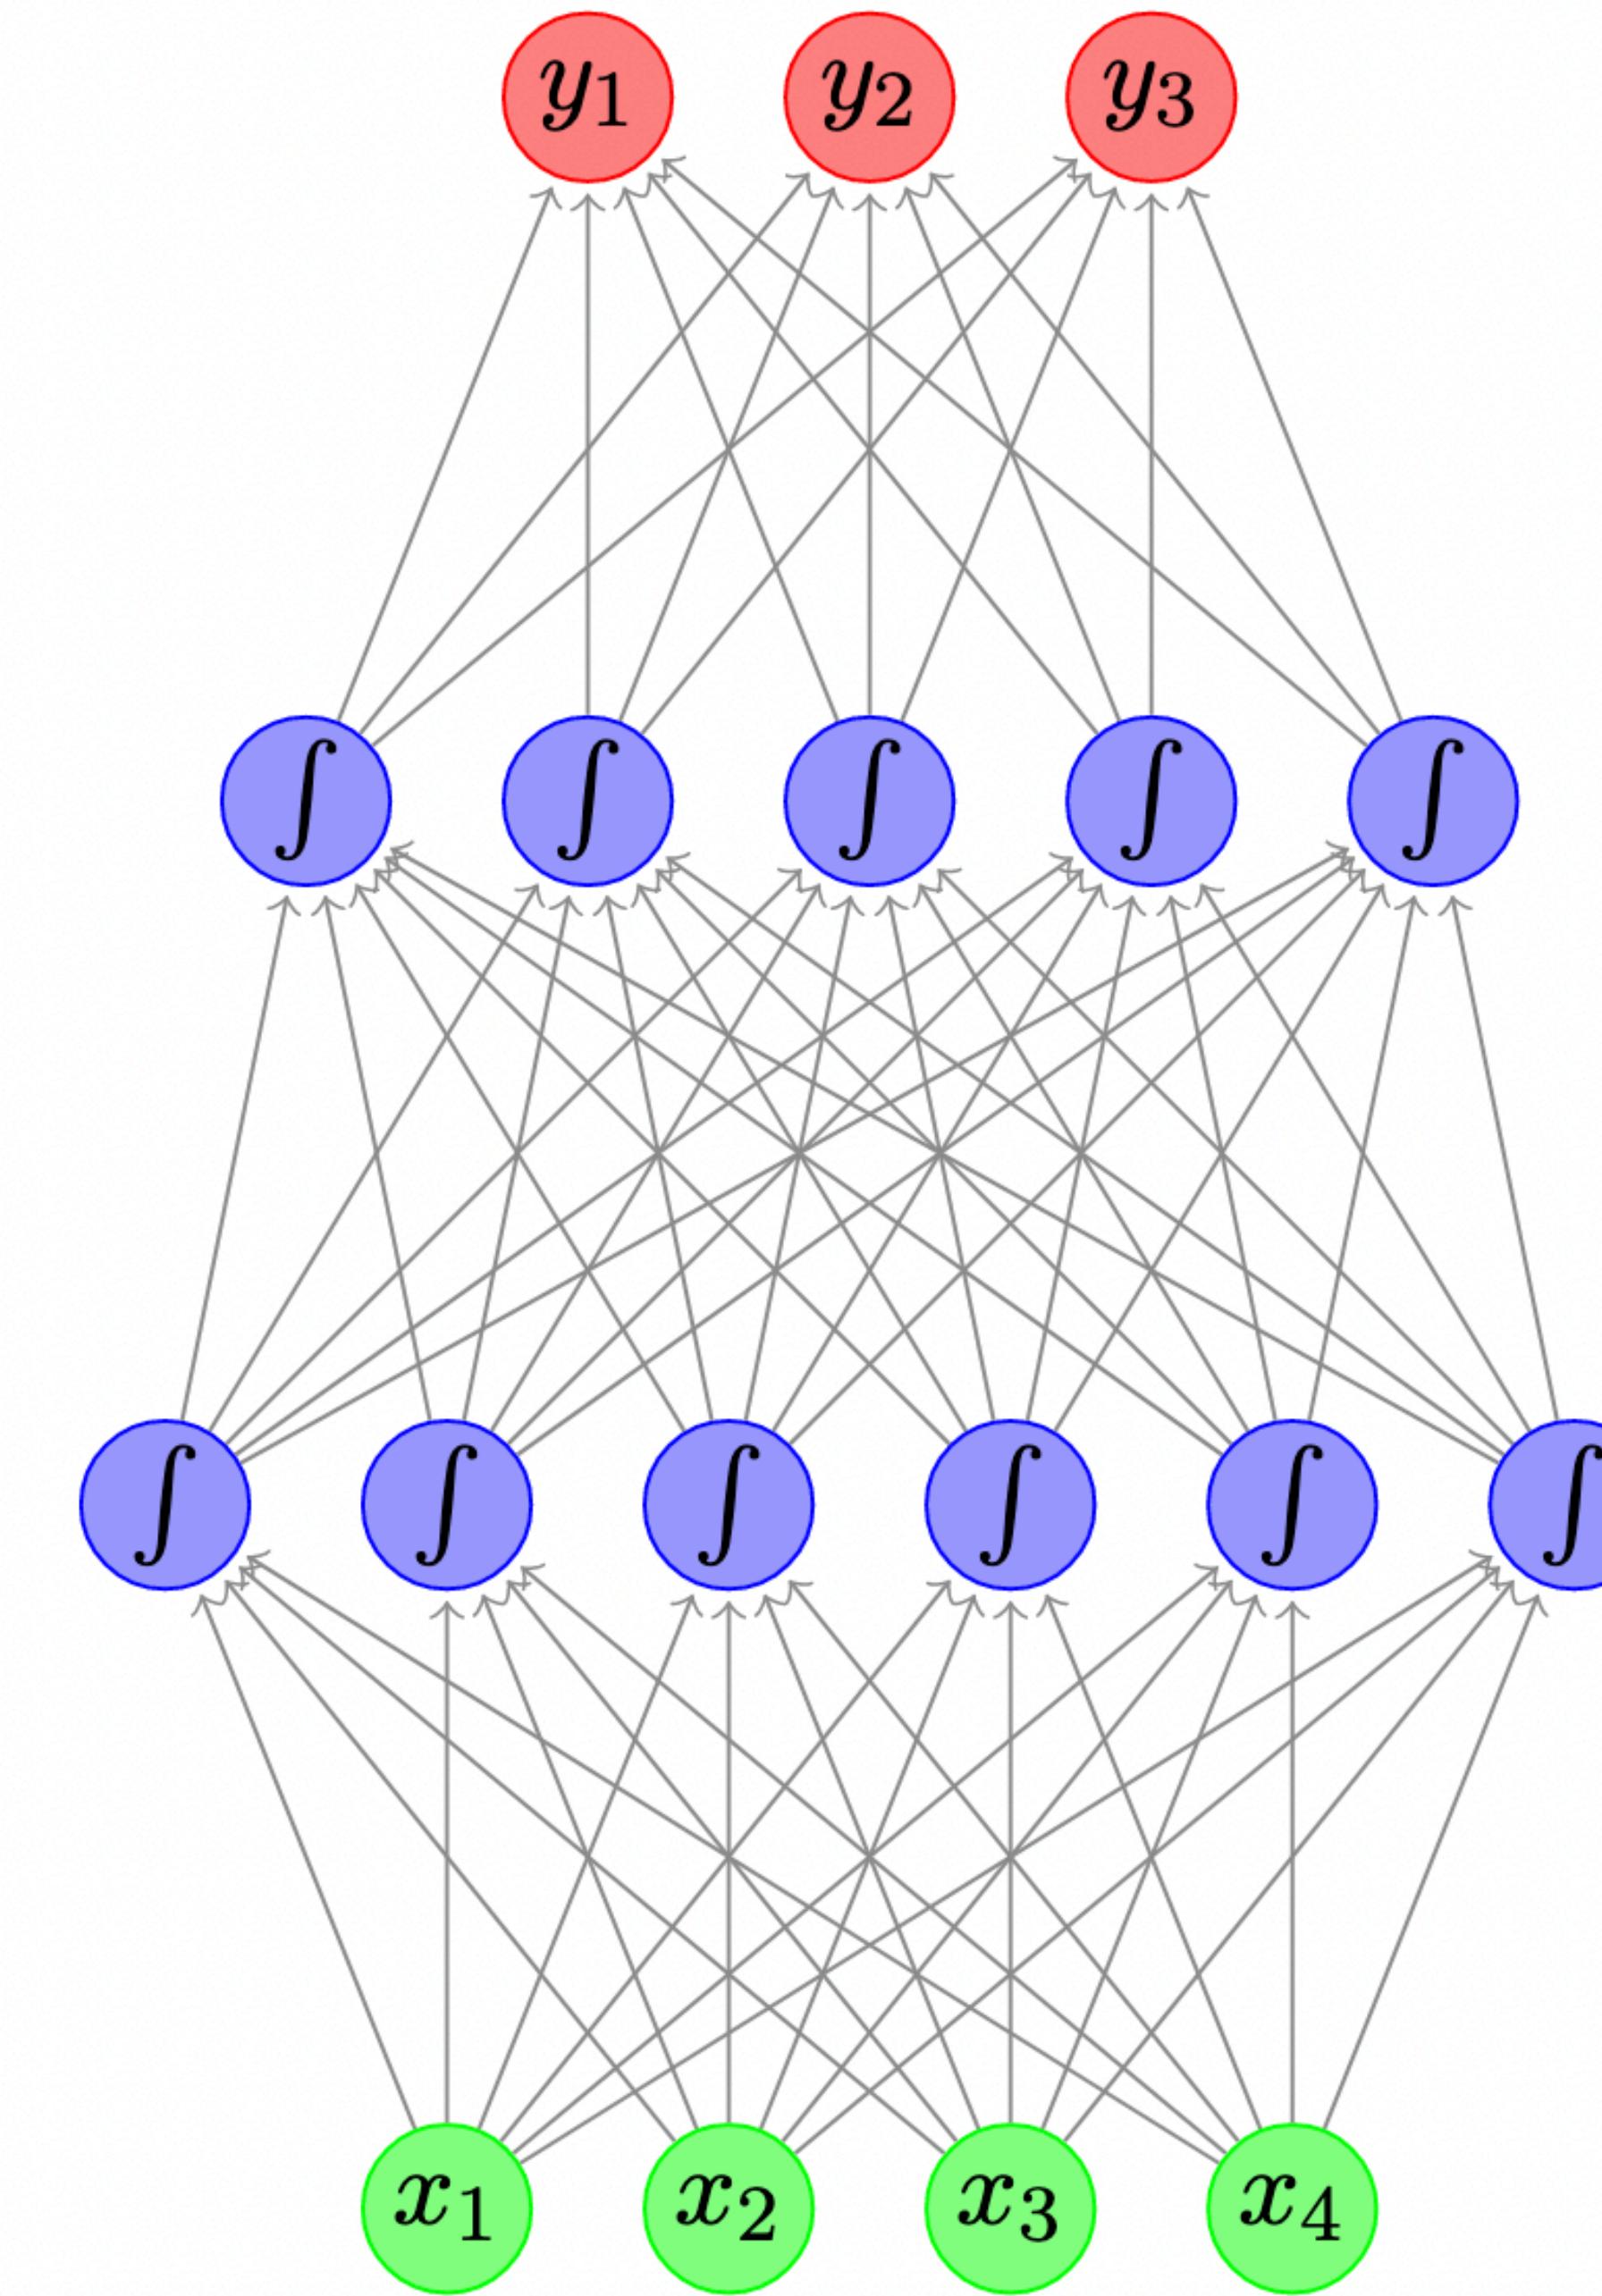
Feed-forward NN

Output
layer

Hidden
layer

Hidden
layer

Input layer



$$g^2(g^1(xW^1 + b^1)W^2 + b^2)$$

$$g^1(xW^1 + b^1)$$

$$x \in \mathbb{R}^{d_{in}}$$

Decisions for the modeller/engineer

1. Number of dimensions

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity
6. Loss function

Decisions for the modeller/engineer

1. Number of dimensions
2. Number of layers
3. Non-linearities (e.g., sigmoid, tanh, rectifier)
4. Output transformation (e.g. softmax)
5. Connectivity
6. Loss function
7. Training regime
(e.g., stochastic gradient descent + flavor;
batching; drop-out)